

1. Download all the data in this folder <https://drive.google.com/open?id=1Z4TyI7FcFVEx8qdl4j09qxvxaqLSqoEu>. it contains two files both images and labels. The label file lists the images and their categories in the following format:

**path/to/the/image.tif,category**

where the categories are numbered 0 to 15, in the following order:

- 0 letter**
- 1 form**
- 2 email**
- 3 handwritten**
- 4 advertisement**
- 5 scientific report**
- 6 scientific publication**
- 7 specification**
- 8 file folder**
- 9 news article**
- 10 budget**
- 11 invoice**
- 12 presentation**
- 13 questionnaire**
- 14 resume**
- 15 memo**

2. On this image data, you have to train 3 types of models as given below. You have to split the data into Train and Validation data.

3. Try not to load all the images into memory, use the generators that we have given the reference notebooks to load the batch of images only during the train data.

or you can use this method also

<https://medium.com/@vijayabhaskar96/tutorial-on-keras-image-datagenerator-with-flow-from-dataframe-8bd5776e45c1>  
(<https://medium.com/@vijayabhaskar96/tutorial-on-keras-image-datagenerator-with-flow-from-dataframe-8bd5776e45c1>)

<https://medium.com/@vijayabhaskar96/tutorial-on-keras-flow-from-dataframe-1fd4493d237c>  
(<https://medium.com/@vijayabhaskar96/tutorial-on-keras-flow-from-dataframe-1fd4493d237c>)

4. You are free to choose Learning rate, optimizer, loss function, image augmentation, any hyperparameters. but you have to use the same architecture what we are asking below.

5. Use tensorboard for every model and analyse your gradients. (you need to upload the screenshots for each model for evaluation)

Note: fit\_generator() method will have problems with the tensorboard histograms, try to debug it, if you could not do use histograms=0 i.e don't include histograms, check the documentation of tensorboard for more information.

6. You can check about Transfer Learning in this link - <https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html>  
(<https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html>)

```
In [1]: !curl --header 'Host: storage.googleapis.com' --user-agent 'Mozilla
```

% Total Time	% Received Current	% Xferd	Average Speed Dload	Speed Upload	Time Total	Time Spent
Left	Speed					
100	4440M	100	4440M	0	0	51.7M
				0	0:01:25	0:01:25
--:--	29.0M					--

```
In [ ]: !unzip data_final.zip -d data_final
```

```
In [3]: %tensorflow_version 2.x
```

```
In [4]: import tensorflow as tf
import os
import numpy as np
import pandas as pd
import warnings
warnings.filterwarnings('ignore')
```

```
In [6]: data = pd.read_csv('labels_final.csv')
data.head()
```

```
Out[6]:
```

	path	label
0	imagesv/v/o/h/voh71d00/509132755+-2755.tif	3
1	imagesl/l/x/t/lxt19d00/502213303.tif	3
2	imagesx/x/e/d/xed05a00/2075325674.tif	2
3	imageso/o/j/b/ojb60d00/517511301+-1301.tif	3
4	imagesq/q/z/k/qzk17e00/2031320195.tif	7

```
In [7]: df = data.groupby('label').tail(2000).reset_index(drop=True)
```

```
In [8]: df.shape
```

```
Out[8]: (32000, 2)
```

```
In [9]: df['label'].value_counts()
```

```
Out[9]:
```

15	2000
14	2000
13	2000
12	2000
11	2000
10	2000
9	2000
8	2000
7	2000
6	2000
5	2000
4	2000
3	2000
2	2000
1	2000
0	2000

Name: label, dtype: int64

```
In [10]: df.head()
```

```
Out[10]:
```

	path	label
0	imagest/t/r/e/tre03e00/2042692622_2042692636.tif	13
1	imagesp/p/n/o/pno92e00/2045613028_2045613030.tif	13
2	imagesa/a/p/d/apd03e00/2042789021_2042789038.tif	13
3	imagesv/v/w/n/vwn30e00/87563496.tif	6
4	imagesu/u/b/w/ubw39c00/2505223727.tif	6

```
In [11]: final = df.sample(frac=1).reset_index(drop=True)
final['label'] = final['label'].apply(str)
```

```
In [12]: final.head()
```

```
Out[12]:
```

	path	label
0	imagesu/u/s/b/usb91a00/1003537721.tif	9
1	imagesf/f/z/x/fzx53a00/1001896119_6122.tif	5
2	imagesa/a/j/a/aja94c00/96383603.tif	0
3	imagesv/v/h/m/vhm79e00/2050743741.tif	13
4	imagesk/k/y/n/kyn20a00/10167827_10167833.tif	6

```
In [13]: from keras.preprocessing.image import ImageDataGenerator
```

```
train = final[:19200]
cv = final[19200:25600]
test = final[25600:]
```

```
In [14]: imageFlow = ImageDataGenerator(rescale=1./255)
```

```
In [15]: train_gen = imageFlow.flow_from_dataframe(train, directory="/content/
                                                classes=None, class_mode=
cv_gen = imageFlow.flow_from_dataframe(cv, directory="/content/data
                                                classes=None, class_mode=
test_gen = imageFlow.flow_from_dataframe(test, directory="/content/
                                                classes=None, class_mode=
```

Found 19200 validated image filenames belonging to 16 classes.  
Found 6400 validated image filenames belonging to 16 classes.  
Found 6400 validated image filenames belonging to 16 classes.

# Transfer Learning

## Model-1

1. Use [VGG-16](https://www.tensorflow.org/api_docs/python/tf/keras/applications/VGG16) ([https://www.tensorflow.org/api\\_docs/python/tf/keras/applications/VGG16](https://www.tensorflow.org/api_docs/python/tf/keras/applications/VGG16)) pretrained network without Fully Connected layers and initialize all the weights with Imagenet trained weights.
2. After VGG-16 network without FC layers, add a new Conv block ( 1 Conv layer and 1 Maxpooling ), 2 FC layers and a output layer to classify 16 classes. You are free to choose any hyperparameters/parameters of conv block, FC layers, output layer.
3. Final architecture will be **INPUT --> VGG-16 without Top layers(FC) --> Conv Layer --> Maxpool Layer --> 2 FC layers --> Output Layer**
4. Train only new Conv block, FC layers, output layer. Don't train the VGG-16 network.

```
In [16]: from keras import applications, callbacks
from tensorflow.keras.layers import Dense, Input, Conv2D, MaxPool2D, Activation
from keras.models import Model
import random as rn
```

```
In [17]: model = applications.VGG16(weights = "imagenet", include_top=False,
model.summary())
```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_notop.h5](https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5) ( [https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_notop.h5](https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5) )  
58892288/58889256 [=====] - 0s 0us/step  
Model: "vgg16"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 156, 256, 3)]	0
block1_conv1 (Conv2D)	(None, 156, 256, 64)	1792
block1_conv2 (Conv2D)	(None, 156, 256, 64)	36928
block1_pool (MaxPooling2D)	(None, 78, 128, 64)	0
block2_conv1 (Conv2D)	(None, 78, 128, 128)	73856
block2_conv2 (Conv2D)	(None, 78, 128, 128)	147584
block2_pool (MaxPooling2D)	(None, 39, 64, 128)	0

block3_conv1 (Conv2D)	(None, 39, 64, 256)	295168
block3_conv2 (Conv2D)	(None, 39, 64, 256)	590080
block3_conv3 (Conv2D)	(None, 39, 64, 256)	590080
block3_pool (MaxPooling2D)	(None, 19, 32, 256)	0
block4_conv1 (Conv2D)	(None, 19, 32, 512)	1180160
block4_conv2 (Conv2D)	(None, 19, 32, 512)	2359808
block4_conv3 (Conv2D)	(None, 19, 32, 512)	2359808
block4_pool (MaxPooling2D)	(None, 9, 16, 512)	0
block5_conv1 (Conv2D)	(None, 9, 16, 512)	2359808
block5_conv2 (Conv2D)	(None, 9, 16, 512)	2359808
block5_conv3 (Conv2D)	(None, 9, 16, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 8, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		

```

In [18]: for layer in model.layers:
          layer.trainable = False
          # print(layer.name)

#adding custom layers
x = model.output

#Conv Layer
Conv1 = Conv2D(filters=32,kernel_size=(3,3),strides=(1,1),padding='
              activation='relu',kernel_initializer=tf.keras.initial
#MaxPool Layer
Pool1 = MaxPool2D(pool_size=(2,2),strides=(2,2),padding='valid',dat

#flatten before adding FC
x = Flatten()(Pool1)

#2 FC layer
fc1 = Dense(256, activation="relu")(x)
fc2 = Dense(256, activation="relu")(fc1)

predictions = Dense(16, activation="softmax")(fc2)

# creating the final model
model_final = Model(model.input, predictions)

# compile the model
model_final.compile(loss = "categorical_crossentropy", optimizer =

```

```

In [19]: model_final.summary()

```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 156, 256, 3)]	0
block1_conv1 (Conv2D)	(None, 156, 256, 64)	1792
block1_conv2 (Conv2D)	(None, 156, 256, 64)	36928
block1_pool (MaxPooling2D)	(None, 78, 128, 64)	0
block2_conv1 (Conv2D)	(None, 78, 128, 128)	73856
block2_conv2 (Conv2D)	(None, 78, 128, 128)	147584
block2_pool (MaxPooling2D)	(None, 39, 64, 128)	0
block3_conv1 (Conv2D)	(None, 39, 64, 256)	295168
block3_conv2 (Conv2D)	(None, 39, 64, 256)	590080

block3_conv3 (Conv2D)	(None, 39, 64, 256)	590080
block3_pool (MaxPooling2D)	(None, 19, 32, 256)	0
block4_conv1 (Conv2D)	(None, 19, 32, 512)	1180160
block4_conv2 (Conv2D)	(None, 19, 32, 512)	2359808
block4_conv3 (Conv2D)	(None, 19, 32, 512)	2359808
block4_pool (MaxPooling2D)	(None, 9, 16, 512)	0
block5_conv1 (Conv2D)	(None, 9, 16, 512)	2359808
block5_conv2 (Conv2D)	(None, 9, 16, 512)	2359808
block5_conv3 (Conv2D)	(None, 9, 16, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 8, 512)	0
Conv1 (Conv2D)	(None, 2, 6, 32)	147488
Pool1 (MaxPooling2D)	(None, 1, 3, 32)	0
flatten (Flatten)	(None, 96)	0
dense (Dense)	(None, 256)	24832
dense_1 (Dense)	(None, 256)	65792
dense_2 (Dense)	(None, 16)	4112
=====		
Total params: 14,956,912		
Trainable params: 242,224		
Non-trainable params: 14,714,688		
=====		

In [20]: `history=model_final.fit_generator(train_gen, steps_per_epoch=600, e`

```
Epoch 1/3
600/600 [=====] - 132s 207ms/step - loss:
1.8984 - accuracy: 0.4009 - val_loss: 1.3349 - val_accuracy: 0.598
0
Epoch 2/3
600/600 [=====] - 97s 162ms/step - loss:
1.2557 - accuracy: 0.6172 - val_loss: 1.2291 - val_accuracy: 0.627
5
Epoch 3/3
600/600 [=====] - 91s 152ms/step - loss:
1.0868 - accuracy: 0.6673 - val_loss: 1.1736 - val_accuracy: 0.645
5
```



```
In [21]: score = model_final.evaluate_generator(test_gen, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

Test score: 1.1859544515609741  
 Test accuracy: 0.6417187452316284

## Model 2

1. Use [VGG-16](https://www.tensorflow.org/api_docs/python/tf/keras/applications/VGG16) ([https://www.tensorflow.org/api\\_docs/python/tf/keras/applications/VGG16](https://www.tensorflow.org/api_docs/python/tf/keras/applications/VGG16)) pretrained network without Fully Connected layers and initialize all the weights with Imagenet trained weights.
2. After VGG-16 network without FC layers, don't use FC layers, use conv layers only as Fully connected layer. any FC layer can be converted to a CONV layer. This conversion will reduce the No of Trainable parameters in FC layers. For example, an FC layer with K=4096 that is looking at some input volume of size 7×7×512 can be equivalently expressed as a CONV layer with F=7,P=0,S=1,K=4096. In other words, we are setting the filter size to be exactly the size of the input volume, and hence the output will simply be 1×1×4096 since only a single depth column “fits” across the input volume, giving identical result as the initial FC layer. You can refer [this \(http://cs231n.github.io/convolutional-networks/#convert\)](http://cs231n.github.io/convolutional-networks/#convert) link to better understanding of using Conv layer in place of fully connected layers.
3. Final architecture will be VGG-16 without FC layers(without top), 2 Conv layers identical to FC layers, 1 output layer for 16 class classification. **INPUT --> VGG-16 without Top layers(FC) --> 2 Conv Layers identical to FC --> Output Layer**
3. Train only last 2 Conv layers identical to FC layers, 1 output layer. Don't train the VGG-16 network.

```
In [22]: imageFlow = ImageDataGenerator(rescale=1./255)
```

```
In [23]: train_gen = imageFlow.flow_from_dataframe(train, directory="/content/
                                                classes=None, class_mode=
cv_gen = imageFlow.flow_from_dataframe(cv, directory="/content/data
                                                classes=None, class_mode=
test_gen = imageFlow.flow_from_dataframe(test, directory="/content/
                                                classes=None, class_mode=
```

Found 19200 validated image filenames belonging to 16 classes.  
Found 6400 validated image filenames belonging to 16 classes.  
Found 6400 validated image filenames belonging to 16 classes.

```
In [24]: model2 = applications.VGG16(weights = "imagenet", include_top=False)
model2.summary()
```

Model: "vgg16"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		

```
In [25]: for layer in model2.layers[:13]:
          # layer.trainable = False
          print(layer.name)
```

```
input_2
block1_conv1
block1_conv2
block1_pool
block2_conv1
block2_conv2
block2_pool
block3_conv1
block3_conv2
block3_conv3
block3_pool
block4_conv1
block4_conv2
```

```
In [26]: for layer in model2.layers:
          layer.trainable = False

#adding custom layers
x = model2.output

#Conv Layers
Conv1 = Conv2D(4096, kernel_size=[7,7], strides=(1,1), padding='val
Conv2 = Conv2D(4096, kernel_size=[1,1], strides=(1,1), padding='val

#flatten before output
x = Flatten()(Conv2)

#output layer
predictions = Dense(16, activation="softmax")(x)

# creating the final model
model2_final = Model(model2.input, predictions)

# compile the model
model2_final.compile(loss = "categorical_crossentropy", optimizer =
```

```
In [27]: model2_final.summary()
```

```
Model: "model_1"
```

Layer (type)	Output Shape	Param #
=====		
input_2 (InputLayer)	[(None, 224, 224, 3)]	0
-----		
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792

block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
conv2d (Conv2D)	(None, 1, 1, 4096)	102764544
conv2d_1 (Conv2D)	(None, 1, 1, 4096)	16781312
flatten_1 (Flatten)	(None, 4096)	0
dense_3 (Dense)	(None, 16)	65552
=====		
Total params: 134,326,096		
Trainable params: 119,611,408		
Non-trainable params: 14,714,688		

```
In [28]: history=model2_final.fit_generator(train_gen, steps_per_epoch=600,
Epoch 1/3
600/600 [=====] - 193s 318ms/step - loss:
2.8943 - accuracy: 0.4394 - val_loss: 1.2508 - val_accuracy: 0.619
1
Epoch 2/3
600/600 [=====] - 191s 318ms/step - loss:
1.0299 - accuracy: 0.6823 - val_loss: 1.0883 - val_accuracy: 0.674
5
Epoch 3/3
600/600 [=====] - 190s 317ms/step - loss:
0.8419 - accuracy: 0.7379 - val_loss: 1.0230 - val_accuracy: 0.692
7
```

```
In [29]: score = model2_final.evaluate_generator(test_gen, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

```
Test score: 1.0471042394638062
Test accuracy: 0.6887500286102295
```

## Model-3

1. Use same network as Model-2 '**INPUT --> VGG-16 without Top layers(FC) --> 2 Conv Layers identical to FC --> Output Layer**' and train only Last 6 Layers of VGG-16 network, 2 Conv layers identical to FC layers, 1 output layer.

```
In [30]: for layer in model2_final.layers[:13]:
        layer.trainable = False
```

```
In [32]: # stopping = callbacks.EarlyStopping(monitor='val_accuracy')
# compile the model
optimizer = tf.keras.optimizers.Adam(learning_rate=0.001, beta_1=0.

model2_final.compile(loss = "categorical_crossentropy", optimizer =
history=model2_final.fit_generator(train_gen, steps_per_epoch=600,
```

Epoch 1/3

600/600 [=====] - 191s 318ms/step - loss: 0.7549 - accuracy: 0.7639 - val\_loss: 1.0614 - val\_accuracy: 0.7031

Epoch 2/3

600/600 [=====] - 190s 317ms/step - loss: 0.6212 - accuracy: 0.8038 - val\_loss: 1.0824 - val\_accuracy: 0.7050

Epoch 3/3

600/600 [=====] - 190s 317ms/step - loss: 0.5538 - accuracy: 0.8255 - val\_loss: 1.0929 - val\_accuracy: 0.7169

```
In [33]: score = model2_final.evaluate_generator(test_gen, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

Test score: 1.126118540763855

Test accuracy: 0.7098437547683716